

Exercise-1

Aim:- Write the installation steps for Hadoop in ubuntu.

Hadoop 2.6.1 installation in ubuntu 16.042-

1) Installing Java

```
$ sudo apt-get update
```

```
$ sudo apt-get install default-jdk
```

it will install java source in your machine at `/usr/lib/jvm/java-7-oracle-jdk-1386`.

To check java version

```
$ java -version
```

2) ADD GROUP & USER

```
$ sudo addgroup hadoop
```

```
$ sudo adduser --ingroup hadoop hduser
```

```
$ sudo adduser hduser sudo
```

```
$ sudo gedit /etc/sudoers
```

in this file at the line "# Allow members of group sudo to execute any command" Add the below line

```
%hduser ALL=(ALL:ALL)ALL
```

LOGOUT the current USER, LOGIN as HDUSER, the below given steps are performed by only HDUSER.

3) Installing SSH

```
$ sudo apt-get install openssh-server
```

Configuring SSH

```
$ sudo su hduser
```

```
$ ssh-keygen
```

```
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized-keys
```

Start ssh.

```
$ sudo /etc/init.d/ssh restart
```

Test your connectivity :

```
$ ssh localhost
```

4) Disabling IPV6

Since hadoop doesn't work on IPV6, we should disable it. One of another reason is also that it has been developed and tested on IPV4 stacks. Hadoop nodes will be able to communicate if we are having IPV4 cluster.

For getting your IPV6 disable in your Linux machine, you need to update /etc/sysctl.conf by

```
$ sudo gedit /etc/sysctl.conf
```

adding following line of codes at end of the file.

```
# disable ipv6
```

```
net.ipv6.conf.all.disable_ipv6 = 1
```

```
net.ipv6.conf.default.disable_ipv6 = 1
```

```
net.ipv6.conf.lo.disable_ipv6 = 1
```

```
$ cat /proc/sys/net/ipv6/conf/all/disable_ipv6
```

it gives "zero" (if system not restart). after

restart the result is "one" (ifus is disable)

Tip:- You can use nano, gedit, and vi editor for updating all text files for this configuration purpose.

5) Download latest Apache Hadoop source from Apache mirrors

* First you need to download Apache Hadoop 2.6.1 (i.e. hadoop-2.6.1.tar.gz) or latest version source from Apache download mirrors. You can also try stable hadoop to get all latest features as well as recent bugs solved with Hadoop source. Choose location where you want to place all your hadoop installation, I have chosen /usr/local/hadoop.

Locate to hadoop installation parent dir

```
$hduser@pingax: ~ $ cd <your hadoop directory>
```

Extract hadoop source

```
$ sudo tar -xvzf hadoop-2.6.1.tar.gz
```

Move hadoop-2.6.1 to hadoop folder

```
$ sudo mv hadoop-2.6.1 /usr/local/hadoop
```

Assign ownership of this folder to Hadoop user

```
$ sudo chown -R hduser:hadoop /usr/local/hadoop
```

```
$ sudo chmod -R 777 /usr/local/hadoop
```

create Hadoop temp directories for namenode and Datanode

```
$ sudo mkdir -p /usr/local/hadoop-tmp/hdfs/namenode
```

```
$ sudo mkdir -p /usr/local/hadoop-tmp/hdfs/datanode
```

Again assign ownership of this Hadoop temp folder to Hadoop user.

```
$ sudo chown -R hduser:hadoop /usr/local/hadoop-tmp/
```

```
$ sudo chmod -R 777 /usr/local/hadoop-tmp/
```

6) update Hadoop Configuration files

User profile : update \$HOME/.bashrc

```
$ sudo gedit ~/.bashrc
```

update hduser configuration file by appending the following environment variables at the end of this file.

```
# -- HADOOP ENVIRONMENT VARIABLES START --
```

```
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-1386
```

```
export PATH=$PATH:$JAVA_HOME/bin
```

```
export HADOOP_HOME=/usr/local/hadoop
```

```
export PATH=$PATH:$HADOOP_HOME/bin
```

```
export PATH=$PATH:$HADOOP_HOME/sbin
```

```
export HADOOP_MAPRED_HOME=$HADOOP_HOME
```

```
export HADOOP_COMMON_HOME=$HADOOP_HOME
```



```
export HADOOP_HDFS_HOME = $HADOOP_HOME
export YARN_HOME = $HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR = $HADOOP_HOME/lib/native
export HADOOP_OPTS = "-Djava.library.path=$HADOOP_HOME/lib"
```

```
# -- HADOOP ENVIRONMENT VARIABLES END -- #
```

Configuration file : hadoop-env.sh.

To edit file, fire the below given command

```
huser@p11ab /usr/local/hadoop/etc/hadoop $ sudo gedit hadoop-env.sh
```

update JAVA_HOME variable,

```
JAVA_HOME = /usr/lib/jvm/java-7-openjdk-1386
```

Configuration file : core-site.xml.

To edit file, fire the below given command

```
/usr/local/hadoop/etc/hadoop $ sudo gedit core-site.xml
```

Paste these lines into <configurations tag.

```
<property>
```

```
<name>fs.default.name</name>
```

```
<value>hdfs://localhost:9000</value>
```

```
</property>
```

Configuration file : hdfs-site.xml.

To edit file, fire the below given command

/usr/local/hadoop/etc/hadoop & sudo gedit hdfs-site.xml

paste these lines into <Configurations> tag.

<property>

<name> dfs.replication </name>

<value> 1 </value>

</property>

<property>

<name> dfs.namenode.name.dir </name>

<value> file:/usr/local/hadoop-tmp/hdfs/
namenode </value>

</property>

<property>

<name> dfs.datanode.data.dir </name>

<value> file:/usr/local/hadoop-tmp/hdfs/
datanode </value>

</property>

Configuration file : yarn-site.xml

To edit file, fire the below given command

/usr/local/hadoop/etc/hadoop & sudo gedit yarn-
site.xml.

paste these lines into <configuration> tag.

<property>

```

<name> yarn, nodemanager, aux-services </name>
<value> mapreduce-shuffle </value>
</property>
<property>
  <name> yarn, nodemanager, aux-services, mapre-
    -duce, shuffle, class </name>
  <value> org.apache.hadoop.mapred.shuffle
    .Handler </value>
</property>

```

Configuration file : mapred-site.xml

Copy template of mapred-site.xml.template file.

```
$ cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.
template
```

```
/usr/local/hadoop/etc/hadoop/mapred-site.xml
```

To edit file, fire the below given command.

```
/usr/local/hadoop/etc/hadoop $ sudo gedit mapred-
-site.xml
```

paste these lines into <configuration> tag

```

<property>
  <name> mapreduce.framework.name </name>
  <value> yarn </value>
</property>

```

Format Namenode

hdfs namenode -format.

Start all Hadoop daemons.

Start hdfs daemons:

`/usr/local/hadoop$ start-dfs.sh.`

Start MapReduce daemons:

`/usr/local/hadoop$ start-yarn.sh.`

Instead of both of these above command you can also use start-all.sh, but its now deprecated so its not recommended to be used for better Hadoop operations.

1) Track/Monitor/Verify.

Verify Hadoop daemons:

`jps.`

In terminal, when we start doing hadoop programs then must and should we have to give the following two commands.

① `start-all.sh.`

② `jps.`

After giving this command, if we get the following 6 nodes, it is successfully installed

JPS

SecondaryNameNode

DataNode

Resource Manager

NodeManager

NameNode

Experiment - 2

Aim :- To implement Word Count Map Reduce Program using standalone hadoop without combiner.

Description:-

* A combiner, also known as a semi-reducer, is an optional class that operates by accepting the inputs from the Map class and thereafter passing the output key-value pairs to the Reducer class.

* The main function of a combiner is to summarize the map output records with the same key.

* The output (key-value collection) of the combiner will be sent over the network to the actual Reducer task as input.

* It is used to reduce the volume of data transfer b/w Map & reduce.

* The input for this program is a text file contains n-lines of text as input.

* The following command is used to run the word count application by taking input files from the input directory.

→ `hadoop jar wc.jar prgmname input_dir output_dir`
to verify resultant files in o/p folder.

→ `hadoop fs -ls output_dir`

to see the output.

→ `hadoop fs -cat output-dir/part-r-00000.`

* The output is in the form of key-value pair as:

`<word, count>`

* ~~not~~

* If we don't use combiner class combiner count of input and output is represented as '0'.

Program:-

WordCount.java:-

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.
    FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.
    FileOutputFormat;

public class WordCount
{
    public static class TokenizerMapper extends
        Mapper<Object, Text, Text, IntWritable>
```

```
private final static IntWritable one = new  
    IntWritable(1);
```

```
private Text word = new Text();
```

```
public void map(Object key, Text value, Context  
    context) throws IOException,  
    InterruptedException
```

```
{  
    StringTokenizer itr = new StringTokenizer  
        (value.toString());
```

```
    while (itr.hasMoreTokens())
```

```
    {  
        word.set(itr.nextToken());
```

```
        context.write(word, one);
```

```
    }
```

```
}
```

```
{  
    public static class IntSumReducer extends  
        Reducer<Text, IntWritable, Text, IntWritable>
```

```
{  
    private IntWritable result = new IntWritable();
```

```
    public void reduce(Text key, Iterable<IntWritable>  
        values, Context context)
```

```
        throws IOException, InterruptedException
```

```
{  
    int sum = 0;
```

```
    for (IntWritable val : values)
```

```
    {  
        sum += val.get();
```

```
    }
```

```

        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws
    Exception
{
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new
        Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(
        args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

a(input Text file):-

This is hadoop lab.
 today is wednesday
 This is important lab
 this is used for analyze big data concepts
 in this we use map reduce

we have to do this

we have to strictly follow the class

this is important lab

This is important lab

this is used for analyze big data concepts

in this we use map reduce

we have to do this

we have to strictly follow the class

this is important lab.

Output:-

This 3

analyze 2

big 2

class 2

concepts 2

data 2

do 2

follow 2

for 2

hadoop 1

have 4

important 4

in 2

is 8

lab 5

map 2

reduce 2

strictly

2

the

2

this

8

to

4

today

1

use

2

used

2

we

6

wednesday 1.

Experiment-3

Aim:- To implement WordCount Map Reduce Program using standalone hadoop with combiner class.

Description:-

- * Combiner class is used in b/w the Map class and the Reduce class to reduce the volume of data transfer b/w map and Reduce.
- * usually, the output of the map task is large and the data transferred to the reduce task is high.
- * combiner class main task is, it uses group by key and groups the $\langle \text{key}, \text{value} \rangle$ pairs of input from map like $(k_1:v_1, k_2:v_2, k_3:v_3)$ to grouping as $(k_1:v_1, v_2, k_2:v_2)$ gives this to Reducer.
- * Its main task is to produce summary information from a large dataset because it replaces the original map output.
- * The input for this program is a text file contains n-lines of text as input (a.txt)
- * The following commands used in execution
hadoop jar $x.jar$, progname ip_dir op_dir
hadoop fs -ls op_dir
hadoop fs -cat op_dir/part-r-00000

* In MapReduce we have RecordReader, RecordWriter.

* RecordWriter ^{takes input} ~~output~~ as <key, value> pair and gives the output as text.

Program:-

Word Count1.

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.
    FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.
    FileOutputFormat;
```

```
public class WordCount1
```

```
{
    public static class TokenizerMapper extends
        Mapper<Object, Text, Text, IntWritable>
```

```
{
    private final static IntWritable one = new
        IntWritable(1);
```

```
    private Text word = new Text();
```

```
    public void
```


public void map (Object Key, Text value, Context
context) throws IOException, InterruptedException

```
{  
    StringTokenizer itr = new StringTokenizer  
        (value.toString());
```

```
    while (itr.hasMoreTokens())
```

```
    {  
        word.set(itr.nextToken());
```

```
        context.write(word, one);
```

```
    }
```

```
}
```

```
}
```

```
public static class IntSumReducer extends  
    Reducer<Text, IntWritable, Text, IntWritable>
```

```
{
```

```
    private IntWritable result = new IntWritable();
```

```
    public void reduce (Text key, Iterable<IntWritable>  
        values, Context context) throws
```

```
        IOException, InterruptedException
```

```
    {
```

```
        int sum = 0;
```

```
        for (IntWritable val : values)
```

```
        {
```

```
            sum += val.get();
```

```
        }
```

```
        result.set(sum);
```

```
        context.write(key, result);
```

```
    }
```

```
}
```

```
public static void main(String[] args) throws  
    Exception
```

```
{  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizeMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path  
        (args[0]));  
    FileOutputFormat.setOutputPath(job, new Path  
        (args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

2

Input:- a.txt. (same as in previous experiment)

Output:-

This	3
analyze	2
big	2
class	2
concepts	2
data	2
do	2
follow	2

for	2
hadoop	1
have	4
important	4
in	2
is	8
lab	5
map	2
reduce	2
strictly	2
the	2
this	8
to	4
today	1
use	2
used	2
we	6
wednesday	1